

Deep learning

Episode 1

Neural networks 101

Maxim Borisyak, Alexander Panin, Andrey Ustyuzhanin



Yandex
Data Factory

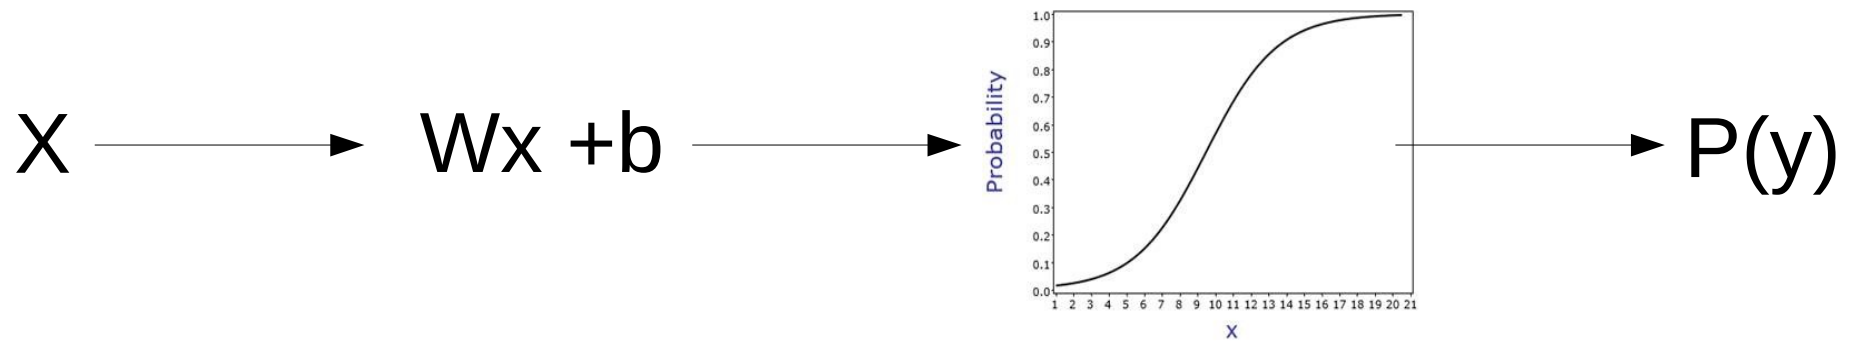
LAMBDA



British Hedgehog
Preservation Society



Recap: logistic regression



Gradient descent

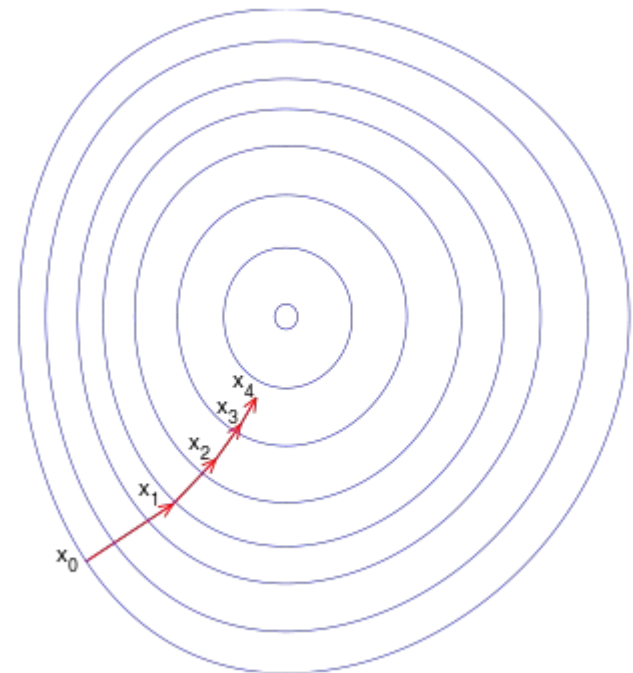
$$P(y|x) = \sigma(w \cdot x + b)$$

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

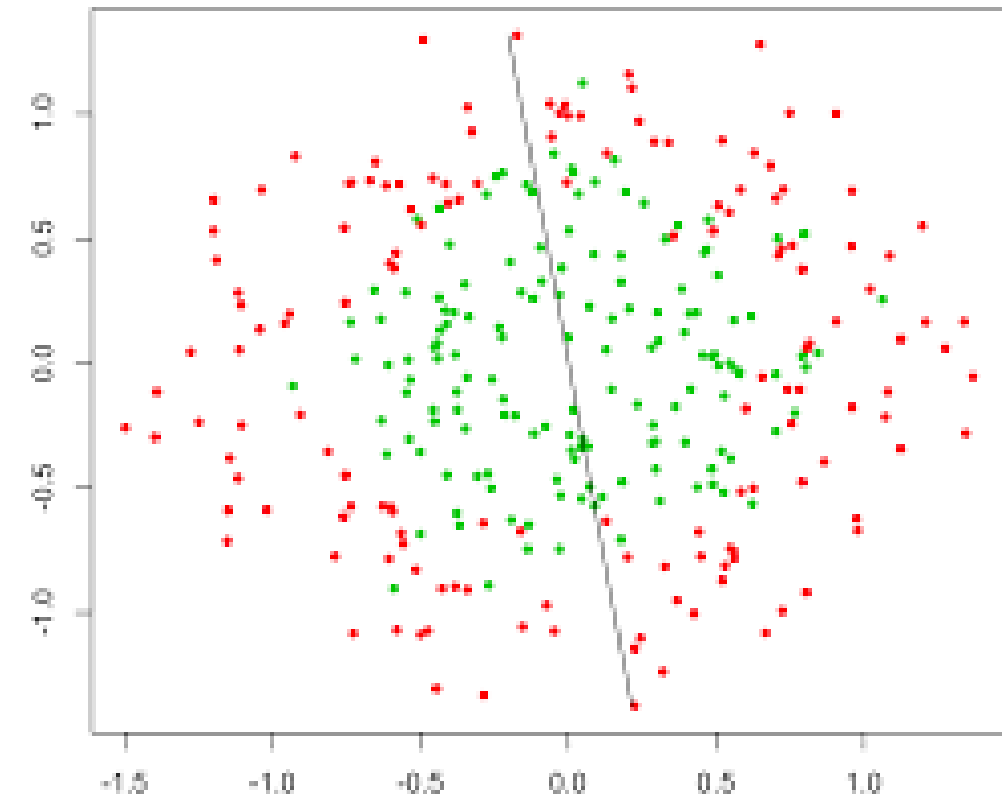
Repeat until convergence

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial L(y, y_{pred})}{\partial \theta_j}$$

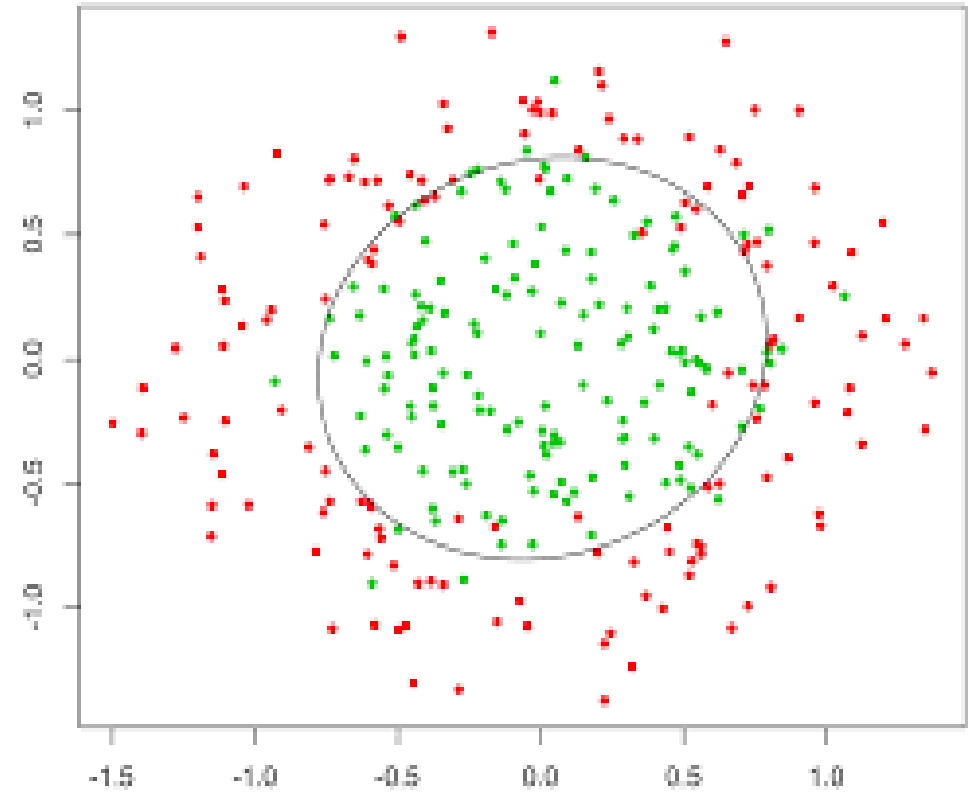
$$\Theta \sim \{w, b\}$$



Nonlinear dependencies



What we have



What we want


- How to get that?

Feature extraction

Loss, for example:


$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:



Training:

$$\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$$




Features would tune to your problem automatically!

What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:




Training: ? $\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$

What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$


Model:



Gradients: $\underset{\theta_2}{\operatorname{argmin}} L(y, P(y|x))$ $\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$

Try linear

Model:




$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

Is it any better than logistic regression?

Try linear


$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

$$w'_i = \sum_j w_j^o w_{ij}^h \qquad b' = \sum_j w_j^o b_j^h + b^o$$

$$P(y|x) = \sigma\left(\sum_i w'_i x_i + b'\right)$$

Try linear

Model:



$$h_j = \sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma \left(\sum_j w_j^o h_j + b^o \right)$$


Output:

$$P(y|x) = \sigma \left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h \right) + b^o \right)$$

Is it any better than logistic regression?

Nonlinearity

Model:

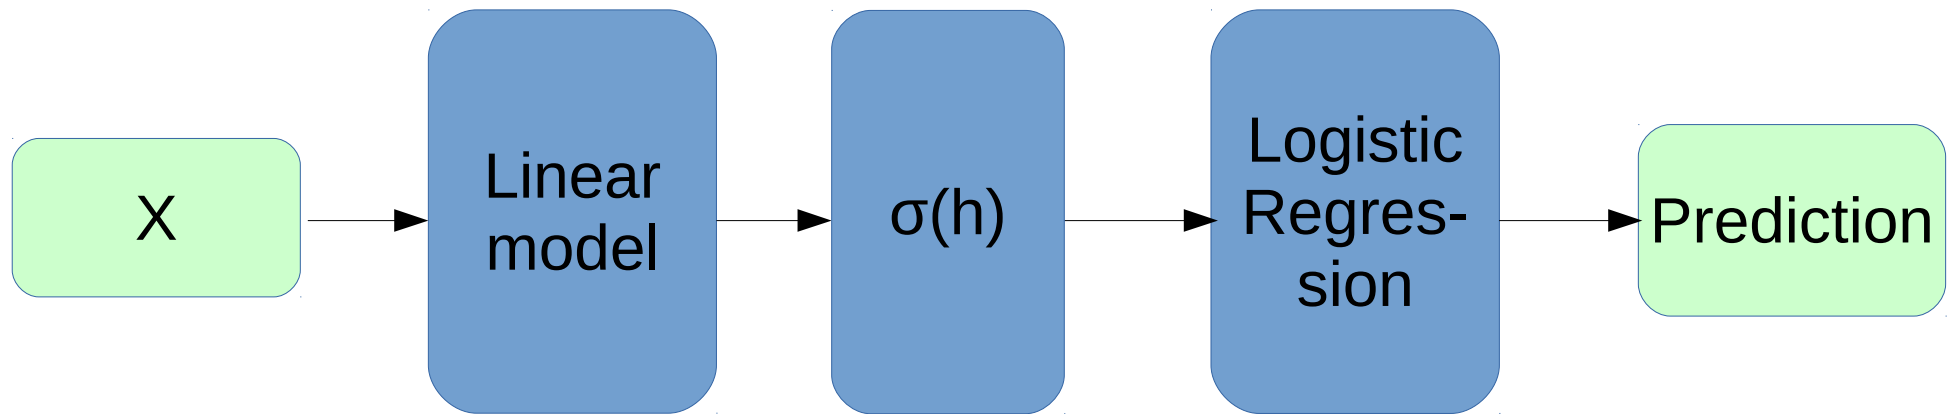


$$h_j = \sigma\left(\sum_{j \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Nonlinearity

Model:



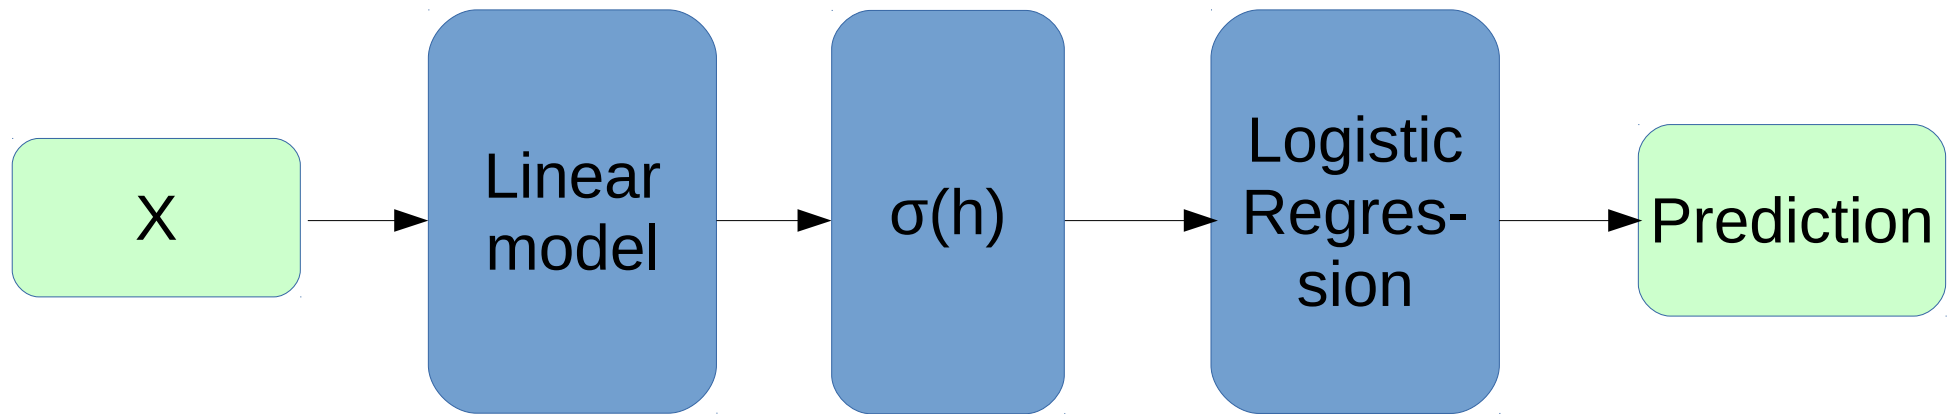
$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right) \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:

$$P(y|x) = \sigma\left(\sum_j w_j^o \sigma\left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$


Nonlinearity

Model:



$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$


$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$



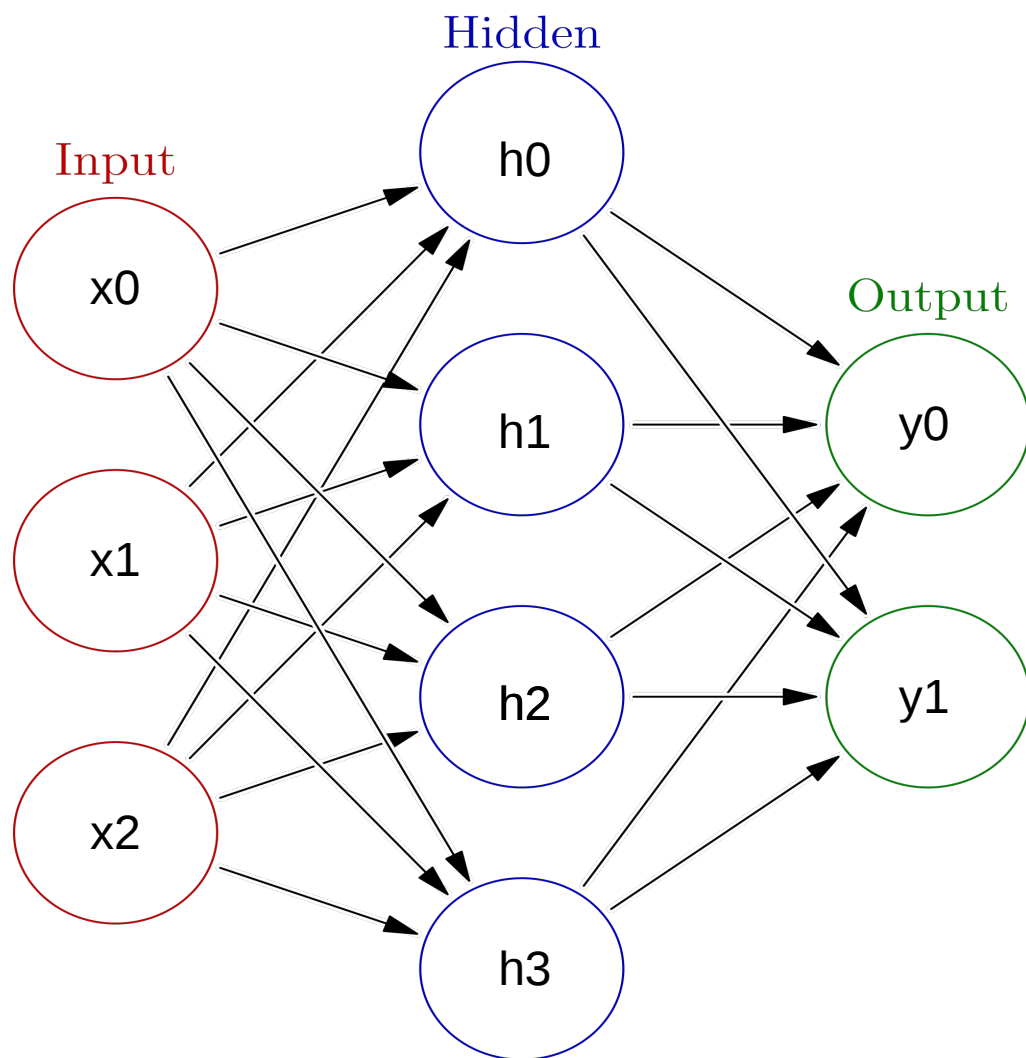
Nonlinearity

- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$

- $f(a) = \max(0, a)$
- $f(a) = \log(1+e^a)$




Initialization, symmetry problem




- Initialize with zeros
 $W \leftarrow 0$
- What will the first step look like?

Initialization, symmetry problem




- Break the symmetry!
- Initialize with random numbers!
 $W \leftarrow N(0, 0.01)?$
 $W \leftarrow U(0, 0.1)?$
- Can get a bit better for deep NNs


Biological inspiration



Biological inspiration




Biological inspiration




Not actual neurons :)

- Neurons react in “spikes”, not real numbers
- Neurons maintain/change their states over time
- No one knows for sure how they “train”
- Neuroglial cells are important
But noone knows, why




Neuroglial Cells of the CNS



Connectionist phrasebook


- Layer – a building block for NNs :
 - “Dense layer”: $f(x) = Wx + b$
 - “Nonlinearity layer”: $f(x) = \sigma(x)$
 - Input layer, output layer
 - A few more we gonna cover later
- Activation – layer output
 - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

Connectionist phrasebook



- “Train it via backprop!”

Connectionist phrasebook




How do we train it?




Discrete Choices


⋮



Layer 2 Features



Layer 1 Features



Original Data


Potential caveats?

Potential caveats?


- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

Regularization

- L1, L2, as usual
- Dropout



(a) Standard Neural Net



(b) After applying dropout.

Computation




The Difference between a CPU and GPU



CPU


MULTIPLE CORES



GPU

THOUSAND OF CORES

Is backprop the only choice?



Nuff

Let's code some neural networks!

